

# Progress Report for CSS 586: Modeling Latent Patterns in Music

Alex Kylo  
akyllo@uw.edu  
University of Washington  
Bothell, WA, USA

## ABSTRACT

This report explores recent research in modeling music with deep learning and provides a progress report on a project to train a generative model of classical music on the MusicNet dataset.

## KEYWORDS

deep learning, neural networks, music

## 1 INTRODUCTION

Machine learning models of music have interesting applications in music information retrieval and creative tools for musical artists and educators. Music is complex and challenging to model because it exhibits a hierarchy of recurring patterns.

Depending on the task, machine learning models of music may be trained on the audio signal itself, either in a time domain or a frequency domain representation, or they may be trained on a digital symbolic representation of music, the most common of which is MIDI (Musical Instrument Digital Interface) notation. MIDI is an encoding of music as streams of bytes in one or more tracks or channels, each representing a sequence of 128 possible pitch values, along with timing, pressure and instrument values. A music transcription model may convert an audio signal into MIDI, which can easily be converted into other symbolic representations such as sheet music, while a synthesizer model can convert MIDI representations into audio signals.

## 2 RELATED WORK

Google's Magenta is an umbrella project for music deep learning research and development of software tools to expose these models for use by creative artists and students.

MusicVAE is a variational LSTM autoencoder for MIDI that incorporates a novel hierarchical structure using a "composer" recurrent layer in its encoder model to better capture structure at multiple levels [5].

Music Transformer is a generative model that borrows its approach from the Natural Language Processing (NLP) domain, using an attention network to model MIDI music as a sequence of discrete tokens with relative positional dependencies [1].

Other research has focused on modeling raw audio waveforms directly. WaveNet is a causal convolutional neural network for generating raw audio waveforms, developed by Google DeepMind, which achieves state of the art performance in generating natural sounding speech from text, but is also capable of generating short, realistic snippets of audio music [4]. Another model named

SampleRNN generates raw audio waveforms using a three-tier hierarchy of gated recurrent units (GRU) to model recurrent structure at multiple temporal resolutions [3].

A paper from Boston University describes an effort to combine the symbolic and waveform approaches to music modeling [2].

## 3 PLANNED METHODS

## 4 PROGRESS

## 5 FUTURE WORK

## REFERENCES

- [1] Cheng-Zhi Anna Huang, Ashish Vaswani, Jakob Uszkoreit, Noam Shazeer, Ian Simon, Curtis Hawthorne, Andrew M. Dai, Matthew D. Hoffman, Monica Dinulescu, and Douglas Eck. 2018. Music Transformer. (2018). arXiv:1809.04281 <http://arxiv.org/abs/1809.04281>
- [2] Rachel Manzelli, Vijay Thakkar, Ali Siahkamari, and Brian Kulis. 2018. Conditioning Deep Generative Raw Audio Models for Structured Automatic Music. *CoRR* abs/1806.09905 (2018). arXiv:1806.09905 <http://arxiv.org/abs/1806.09905>
- [3] Soroush Mehri, Kundan Kumar, Ishaan Gulrajani, Rithesh Kumar, Shubham Jain, Jose Sotelo, Aaron Courville, and Yoshua Bengio. [n.d.]. SampleRNN: An Unconditional End-to-End Neural Audio Generation Model. ([n. d.]). arXiv:1612.07837 <http://arxiv.org/abs/1612.07837>
- [4] Aaron van den Oord, Sander Dieleman, Heiga Zen, Karen Simonyan, Oriol Vinyals, Alex Graves, Nal Kalchbrenner, Andrew Senior, and Koray Kavukcuoglu. [n.d.]. WaveNet: A Generative Model for Raw Audio. ([n. d.]). arXiv:1609.03499 <http://arxiv.org/abs/1609.03499>
- [5] Adam Roberts, Jesse Engel, Colin Raffel, Curtis Hawthorne, and Douglas Eck. [n.d.]. A Hierarchical Latent Vector Model for Learning Long-Term Structure in Music. In *Proceedings of the 35th International Conference on Machine Learning* (2018). 10.