

Progress Report for CSS 586 Project: Music Generation Using Deep Learning

Jack Phan
phan92@uw.edu
University of Washington
Bothell, WA, USA

ABSTRACT

The original goal of the project is music transcription, which is the translation of audio data to symbolic representation of musical notations. With further consideration, we decided to focus on a different problem to better align our work. The new objective is music generation using deep learning. There are several methods that could be used to generate new music from a corpus of training data: MIDI, time-domain sound pressure data, frequency-domain representation of the sound pressure. Regardless of the data type, music generation involves the application of a sequence model such as recurrent neural network (RNN). The focus of my contribution is using a combination of time-domain and frequency-domain information. The current progress includes literature review, data exploration, learning about different representations of an audio signal. We are working on building a data pipeline and experimenting with various RNN models.

KEYWORDS

deep learning, neural networks, music generation, RNN

1 INTRODUCTION

Recurrent Neural Network (RNN) has shown its capability in recent years to model time-series sequences. Advanced version of RNN such as Gated Recurrent Unit (GRU) and Long Short-Term Memory (LSTM) have been successfully used for applications such as voice recognition, text to speech, and other applications that involve time-series data. RNN could also be used to generate novel sequences of texts or sounds.

A music recording is a digitized time-varying sequence of pressure displacement. Each point in the sequence represents the amplitude of the pressure strength at a moment in time. The sequence could also be converted to the frequency-domain representation, which describes its sound components in terms of amplitude and frequency. A popular representation of a sound segment is spectrogram. The goal of this project is generative modeling for short polyphonic music segments. We will build and train deep learning models using spectrogram as input, and generate novel music segments using these models.

2 RELATED WORK

Musical generative modeling is an interesting area of research. Recent innovations in this domain involve deep learning using a variety of modelling techniques. This section presents some of the recent works in this domain.

Historically, the modeling of raw audio data is extremely challenging because a single second of audio recording could contains

up to 44,100 samples. Google DeepMind's WaveNet is one of the first models that successfully learn from raw audio data [1]. It is designed using a technique called 1D dilated causal convolution that allows an output to capture information from many inputs with a minimal computational cost. WaveNet was designed for the primary purpose of generating speech which mimics any human voice. Since the architecture can be used to model any raw audio signal, it was also used to generate music. Although the outputs were far from being any masterpieces, the results show the possibility of generating musical pieces from the raw audio data.

An alternative representation of the time-series signal is the frequency-domain representation. Existing generative models for audio have largely focused on the time-domain waveforms. In 2019, a generative model for audio in the frequency domain called MelNet was introduced [3]. MelNet models the spectrograms, which are time-frequency representations of audio. The advantage of spectrograms is that they are compressed expressions of the time-domain signal. MelNet leverages this advantage to generate high-fidelity audio samples, which capture structures that are still challenging for time-domain models.

3 PLANNED WORK

The focus of this project is generating short musical pieces by modeling the spectrograms of input audio. The following subsections describe our plan toward accomplishing this objective.

3.1 Datasets

We use a dataset called MusicNet published by University of Washington, the department of Computer Science [2]. This dataset is a collection of 330 freely-licensed classical music recordings captured under various studio and microphone conditions. The recordings are in .wav raw audio format and are polyphonic, meaning that multiple musical notes could present at a single time point. Furthermore, the music is performed by multiple instruments, which have unique sound signatures.

3.2 Data Processing

Each training example is a raw sound fragment, which has a length of several seconds, taken from an audio recording. We will transform the time-domain sample to the mel spectrogram representation. The mel scale is a logarithmic transformation of a signal frequency. The idea of this transformation is to replicate the human perception of the difference between two sounds of different frequencies. For example, a person can easily tell the difference between 100 Hz and 200 Hz, but cannot distinguish 1000 Hz and 1100 Hz. This transformation is done by splitting the sample into

overlapping windows, performing the Short Time Fourier Transformation on each window, converting the frequencies to mel scale, and finally converting the amplitude to decibels. Librosa, a python package for music analysis, has built-in functions for performing this transformation. An example of a mel spectrogram is shown in Figure 1.




Figure 1: Mel spectrogram of a 10 second sound segment from the dataset. The color represents the amplitude in dB.

3.3 Model Building

We will implement sequence models such as RNN, GRU, and LSTM to model the mel spectrogram and use these models to generate novel spectrograms that could be converted back to audio.

4 PROGRESS

Initially, a lot of time was spent on understanding different representation of the input data. We also learned about the basics of sequence models such as RNN, GRU, and LSTM. Because the dataset is large and cannot be loaded into the memory at once, we are working on a data pipeline that processes the raw audio files as described in the Section Data Processing and presents training examples one by one to a model.

REFERENCES

- [1] Aaron van den Oord, Sander Dieleman, Heiga Zen, Karen Simonyan, Oriol Vinyals, Alex Graves, Nal Kalchbrenner, Andrew Senior, and Koray Kavukcuoglu. 2016. WaveNet: A Generative Model for Raw Audio. (2016). arXiv:1609.03499 <http://arxiv.org/abs/1609.03499>
- [2] John Thickstun, Zaid Harchaoui, and Sham Kakade. 2017. Learning Features of Music from Scratch. arXiv:stat.ML/1611.09827
- [3] Sean Vasequez and Mike Lewis. 2019. MelNet: A Generative Model for Audio in the Frequency Domain. arXiv:eess.AS/1906.01083